

Sai Kiran Reddy Kotha

Banner 700746206

Link:

https://drive.google.com/file/d/1qe828APPGYQwurC_EK16CMR5R5XGZwMf/view?usp=share_link

1. Read the provided CSV file 'data.csv'.

<https://drive.google.com/drive/folders/1h8C3mLsso-R-sIOLsvoYwPLzy2fJ4IOF?usp=sharing>

```
In [81]: import pandas as pd
import numpy as np
import random as rnd

# visualization
import seaborn as sns
import matplotlib.pyplot as plt

# machine learning
from sklearn.linear_model import LogisticRegression, RidgeClassifierCV
from sklearn.svm import SVC, LinearSVC
from sklearn.ensemble import (RandomForestClassifier, GradientBoostingClassifier)
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import Perceptron
from sklearn.linear_model import SGDClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import cross_val_score, GridSearchCV
from sklearn.metrics import accuracy_score

from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, recall_score, precision_score, classification_report, confusion_matrix
import warnings # current version generates a bunch of warnings that we'll ignore
warnings.filterwarnings("ignore")

In [82]: df = pd.read_csv(r"C:\Users\ksaik\Downloads\data.csv")
print(df.head())
```

	Duration	Pulse	Maxpulse	Calories
0	60	110	130	409.1
1	60	117	145	479.0
2	60	103	135	340.0
3	45	109	175	282.4
4	45	117	148	406.0

2. Show the basic statistical description about the data.
3. Check if the data has null values.
 - a. Replace the null values with the mean

```
In [83]: print(df.describe())
```

	Duration	Pulse	Maxpulse	Calories
count	169.000000	169.000000	169.000000	164.000000
mean	63.846154	107.461538	134.047337	375.790244
std	42.299949	14.510259	16.450434	266.379919
min	15.000000	80.000000	100.000000	50.300000
25%	45.000000	100.000000	124.000000	250.925000
50%	60.000000	105.000000	131.000000	318.600000
75%	60.000000	111.000000	141.000000	387.600000
max	300.000000	159.000000	184.000000	1860.400000

```
In [84]: df.isnull().any()
```

```
Out[84]: Duration      False
Pulse                False
Maxpulse             False
Calories              True
dtype: bool
```

```
In [85]: df.fillna(df.mean(), inplace=True)
```

4. Select at least two columns and aggregate the data using: min, max, count, mean.

```
In [86]: df.agg({'Duration': ['min', 'max', 'count', 'mean'], 'Pulse': ['min', 'max', 'count', 'mean']})
```

Out[86]:

	Duration	Pulse
min	15.000000	80.000000
max	300.000000	159.000000
count	169.000000	169.000000
mean	63.846154	107.461538

5. Filter the dataframe to select the rows with calories values between 500 and 1000.

```
In [87]: df.loc[(df['Calories']>500)&(df['Calories']<1000)]
```

Out[87]:

	Duration	Pulse	Maxpulse	Calories
51	80	123	146	643.1
62	160	109	135	853.0
65	180	90	130	800.4
66	150	105	135	873.4
67	150	107	130	816.0
72	90	100	127	700.0
73	150	97	127	953.2
75	90	98	125	563.2
78	120	100	130	500.4
90	180	101	127	600.1
99	90	93	124	604.1
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3

6. Filter the dataframe to select the rows with calories values > 500 and pulse < 100.

```
In [88]: df.loc[(df['Calories']>500)&(df['Pulse']<100)]
```

Out[88]:

	Duration	Pulse	Maxpulse	Calories
65	180	90	130	800.4
70	150	97	129	1115.0
73	150	97	127	953.2
75	90	98	125	563.2
99	90	93	124	604.1
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3

7. Create a new “df_modified” dataframe that contains all the columns from df except for “Maxpulse”.

```
In [89]: df_modified = df[['Duration', 'Pulse', 'Calories']]
df_modified.head()
```

Out[89]:

	Duration	Pulse	Calories
0	60	110	409.1
1	60	117	479.0
2	60	103	340.0
3	45	109	282.4
4	45	117	406.0

8. Delete the “Maxpulse” column from the main df dataframe

```
In [90]: del df['Maxpulse']
```

```
In [91]: df.head()
```

Out[91]:

	Duration	Pulse	Calories
0	60	110	409.1
1	60	117	479.0
2	60	103	340.0
3	45	109	282.4
4	45	117	406.0

9. Convert the datatype of Calories column to int datatype.

```
In [92]: df.dtypes
```


```
Out[92]: Duration      int64  
Pulse      int64  
Calories    float64  
dtype: object
```

```
In [93]: df['Calories'] = df['Calories'].astype(np.int64)  
df.dtypes
```

```
Out[93]: Duration      int64  
Pulse      int64  
Calories    int64  
dtype: object
```


10. Using pandas create a scatter plot for the two columns (Duration and Calories).

a. Example:



```
In [94]: df.plot.scatter(x='Duration',y='Calories',c='DarkBlue')
```

```
Out[94]: <AxesSubplot:xlabel='Duration', ylabel='Calories'>
```



1. (Titanic Dataset)

- Find the correlation between 'survived' (target column) and 'sex' column for the Titanic use case in class.
 - Do you think we should keep this feature?

```
In [97]: df=pd.read_csv(r"C:\Users\ksaik\Downloads\Dataset\Dataset\train.csv")
df.head()
```

```
Out[97]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

```
In [98]: #converted categorical data to numerical values for correlation calculation
```

```
label_encoder = preprocessing.LabelEncoder()
df['Sex'] = label_encoder.fit_transform(df.Sex.values)


#Calculation of correlation for 'Survived' and 'Sex' in data
correlation_Value= df['Survived'].corr(df['Sex'])

print(correlation_Value)

-0.543351380657755
```

2. Do at least two visualizations to describe or show correlations.

```
In [99]: #2 2. Do at Least two visualizations to describe or show correlations.
# using heatmap from seaborn
corr=df.corr()
sns.heatmap(corr, vmax=1, vmin=-1, center=0, cmap='vlag')
plt.show()
```



```
In [100]: df.corr().style.background_gradient()
```

```
Out[100]:
```

	PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Fare
PassengerId	1.000000	-0.005007	-0.035144	0.042939	0.036847	-0.057527	-0.001652	0.012658
Survived	-0.005007	1.000000	-0.338481	-0.543351	-0.077221	-0.035322	0.081629	0.257307
Pclass	-0.035144	-0.338481	1.000000	0.131900	-0.369226	0.083081	0.018443	-0.549500
Sex	0.042939	-0.543351	0.131900	1.000000	0.093254	-0.114631	-0.245489	-0.182333
Age	0.036847	-0.077221	-0.369226	0.093254	1.000000	-0.308247	-0.189119	0.096067
SibSp	-0.057527	-0.035322	0.083081	-0.114631	-0.308247	1.000000	0.414838	0.159651
Parch	-0.001652	0.081629	0.018443	-0.245489	-0.189119	0.414838	1.000000	0.216225
Fare	0.012658	0.257307	-0.549500	-0.182333	0.096067	0.159651	0.216225	1.000000

3. Implement Naïve Bayes method using scikit-learn library and report the accuracy.

```
In [101]: #3. Implement Naïve Bayes method using scikit-learn library and report the accuracy.
```

```
train_raw = pd.read_csv(r"C:\Users\ksaik\Downloads\Dataset\Dataset\train.csv")
test_raw = pd.read_csv(r"C:\Users\ksaik\Downloads\Dataset\Dataset\test.csv")
train_raw['train'] = 1
test_raw['train'] = 0
df = train_raw.append(test_raw, sort=False)
features = ['Age', 'Embarked', 'Fare', 'Parch', 'Pclass', 'Sex', 'SibSp']
target = 'Survived'
df = df[features + [target] + ['train']]
df['Sex'] = df['Sex'].replace(["female", "male"], [0, 1])
df['Embarked'] = df['Embarked'].replace(['S', 'C', 'Q'], [1, 2, 3])
train = df.query('train == 1')
test = df.query('train == 0')
train.dropna(axis=0, inplace=True)
labels = train[target].values
train.drop(['train', target, 'Pclass'], axis=1, inplace=True)
test.drop(['train', target, 'Pclass'], axis=1, inplace=True)
```

```
In [102]: #Test and train split
X_train, X_val, Y_train, Y_val = train_test_split(train, labels, test_size=0.25, random_state=1)
classifier = GaussianNB()
classifier.fit(X_train, Y_train)
y_pred = classifier.predict(X_val)
```

```
# Summary of the predictions made by the classifier
```

```
print(classification_report(Y_val, y_pred))
print(confusion_matrix(Y_val, y_pred))
# Accuracy score
```

```
print('accuracy is', accuracy_score(Y_val, y_pred))
```

```

              precision    recall  f1-score   support

    0.0         0.80      0.81      0.81         102
    1.0         0.74      0.72      0.73          76

 accuracy          0.77
 macro avg         0.77      0.77      0.77         178
weighted avg         0.77      0.78      0.77         178
```

```
[[83 19]
 [21 55]]
accuracy is 0.7752808988764045
```

2. (Glass Dataset)

1. Implement Naïve Bayes method using scikit-learn library.
 - a. Use the glass dataset available in [Link](#) also provided in your assignment.
 - b. Use **train_test_split** to create training and testing part.
2. Evaluate the model on testing part using score and

```
classification_report(y_true, y_pred)
```

```
In [103]: """(Glass Dataset)
1. Implement Naïve Bayes method using scikit-learn library.
a. Use the glass dataset available in Link also provided in your assignment.
b. Use train_test_split to create training and testing part.
2. Evaluate the model on testing part using score and"""

# reading the dataset
glass=pd.read_csv(r"C:\Users\ksaik\Downloads\Dataset\Dataset\glass.csv")

features = ['Rl', 'Na', 'Mg', 'Al', 'Si', 'K', 'Ca', 'Ba', 'Fe']
target = 'Type'
X_train, X_val, Y_train, Y_val = train_test_split(glass[:, :-1], glass['Type'], test_size=0.25, random_state=1)
classifier = GaussianNB()
classifier.fit(X_train, Y_train)
y_pred = classifier.predict(X_val)
# Summary of the predictions made by the classifier
print(classification_report(Y_val, y_pred))
print(confusion_matrix(Y_val, y_pred))
# Accuracy score

print('accuracy is', accuracy_score(Y_val, y_pred))
```

	precision	recall	f1-score	support
1	0.92	0.96	0.94	23
2	0.94	0.94	0.94	16
3	0.75	0.43	0.55	7
5	0.00	0.00	0.00	1
6	1.00	1.00	1.00	1
7	0.67	0.67	0.67	6
accuracy			0.83	54
macro avg	0.71	0.66	0.68	54
weighted avg	0.86	0.83	0.84	54

```
[[22 1 0 0 0 0]
 [ 1 15 0 0 0 0]
 [ 1 0 3 2 0 1]
 [ 0 0 0 0 0 1]
 [ 0 0 0 0 1 0]
 [ 0 0 1 1 0 4]]
accuracy is 0.8333333333333334
```

1. Implement linear SVM method using scikit library
 - a. Use the glass dataset available in [Link](#) also provided in your assignment.
 - b. Use **train_test_split** to create training and testing part.
2. Evaluate the model on testing part using score and

```
classification_report(y_true, y_pred)
```

```
In [104]: classifier = LinearSVC()

classifier.fit(X_train, Y_train)

y_pred = classifier.predict(X_val)

# Summary of the predictions made by the classifier
print(classification_report(Y_val, y_pred))
print(confusion_matrix(Y_val, y_pred))
print('accuracy is', accuracy_score(Y_val, y_pred))
```

```

              precision    recall  f1-score   support

     1         1.00      0.91      0.95         23
     2         0.48      1.00      0.65         16
     3         0.00      0.00      0.00          7
     5         0.00      0.00      0.00          1
     6         0.00      0.00      0.00          1
     7         0.00      0.00      0.00          6

 accuracy          0.69         54
 macro avg         0.25         0.32         0.27         54
 weighted avg      0.57         0.69         0.60         54

[[21  2  0  0  0  0]
 [ 0 16  0  0  0  0]
 [ 0  7  0  0  0  0]
 [ 0  1  0  0  0  0]
 [ 0  1  0  0  0  0]
 [ 0  6  0  0  0  0]]
accuracy is 0.6851851851851852
```

Do at least two visualizations to describe or show correlations in the Glass Dataset.

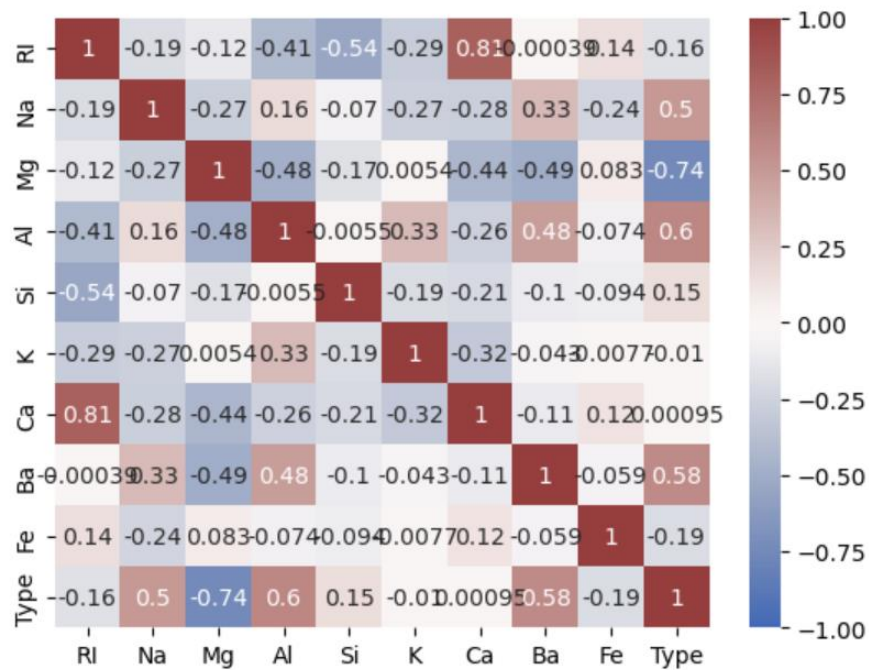
```
In [105]: glass.corr().style.background_gradient()
```

Out[105]:

	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	Type
RI	1.000000	-0.191885	-0.122274	-0.407326	-0.542052	-0.289833	0.810403	-0.000386	0.143010	-0.164237
Na	-0.191885	1.000000	-0.273732	0.156794	-0.069809	-0.266087	-0.275442	0.326603	-0.241346	0.502898
Mg	-0.122274	-0.273732	1.000000	-0.481799	-0.165927	0.005396	-0.443750	-0.492262	0.083060	-0.744993
Al	-0.407326	0.156794	-0.481799	1.000000	-0.005524	0.325958	-0.259592	0.479404	-0.074402	0.598829
Si	-0.542052	-0.069809	-0.165927	-0.005524	1.000000	-0.193331	-0.208732	-0.102151	-0.094201	0.151565
K	-0.289833	-0.266087	0.005396	0.325958	-0.193331	1.000000	-0.317836	-0.042618	-0.007719	-0.010054
Ca	0.810403	-0.275442	-0.443750	-0.259592	-0.208732	-0.317836	1.000000	-0.112841	0.124968	0.000952
Ba	-0.000386	0.326603	-0.492262	0.479404	-0.102151	-0.042618	-0.112841	1.000000	-0.058692	0.575161
Fe	0.143010	-0.241346	0.083060	-0.074402	-0.094201	-0.007719	0.124968	-0.058692	1.000000	-0.188278
Type	-0.164237	0.502898	-0.744993	0.598829	0.151565	-0.010054	0.000952	0.575161	-0.188278	1.000000

In [106]:

```
sns.heatmap(glass.corr(), annot=True, vmax=1, vmin=-1, center=0, cmap='vlag')
plt.show()
```



Which algorithm you got better accuracy? Can you justify why?

We can numerically conclude that Naive bayes is better than SVM because there is less correlation in data and the accuracy is also higher